

AUDITABLE ETL PIPELINES IN REGULATED ENVIRONMENTS USING A DEVSECOPS-BASED TOOLCHAIN

Daniel Dickerson

Bachelor of Engineering (Honours)
Software Engineering



School of Engineering
Macquarie University

November 2025

Supervisor: Dr Ansgar Fehnker

ACKNOWLEDGMENTS

I would like to thank my supervisor, Dr Ansgar Fehnker, for his guidance and feedback throughout this project; Dr Kate Stefanov, for her ongoing advice and coordination through the weekly thesis sessions; and my industry co-supervisors, Aneesh Kondepuri and Richard Hawkes, for providing feedback and assistance throughout this project.

STATEMENT OF CANDIDATE

I, Daniel Dickerson, declare that this report, submitted as part of the requirement for the award of Bachelor of Engineering in the School of Engineering, Macquarie University, is entirely my own work unless otherwise referenced or acknowledged. This document has not been submitted for qualification or assessment at any other academic institution.

Student's Name: Daniel Dickerson

Student's Signature: Daniel Dickerson

Date: May 20, 2026

ABSTRACT

This thesis explores how a DevSecOps toolchain can be designed to support ETL (Extract–Transform–Load) workflows in regulated industries, with a focus on auditability and governance. It applies a Design Science Research approach to design a control framework that translates regulatory expectations—such as those in APRA’s CPS 230, 234, and 510—into technical assurance mechanisms.

The work to date defines a conceptual baseline pipeline using Apache Airflow for orchestration, PySpark on EMR for transformation, and Amazon S3 for storage. Infrastructure is provisioned through Terraform and deployed via GitHub Actions, forming a reproducible but ungoverned model that establishes how data moves, where control resides, and what evidence is naturally produced. Analysis of this baseline identifies gaps in governance, integrity verification, and end-to-end lineage, forming the foundation for the DevSecOps architecture that follows.

The study positions DevSecOps as a bridge between software assurance and data governance, proposing that automation, Infrastructure-as-Code, and policy enforcement can transform ETL pipelines from operational utilities into verifiable systems of record. The research completed so far defines the regulatory context, design requirements, and evaluation framework that will guide the subsequent construction and assessment of the DevSecOps-enabled pipeline.

Contents

Acknowledgments	ii
Abstract	iv
Table of Contents	v
List of Figures	vii
List of Tables	viii
1 Introduction	1
1.1 Context and Motivation	1
1.2 Problem Statement	2
1.3 Research Objectives	2
1.4 Contributions	3
1.5 Project Timeline	3
1.6 Thesis Structure	3
2 Literature Review	5
2.1 Background	5
2.2 Defining Key Concepts	6
2.3 RQ1: ETL Pipelines in Regulated Industries	9
2.4 RQ2: DevSecOps and Governance in the SDLC	10
2.5 RQ3: Comparing the Data Lifecycle and the SDLC	11
2.6 RQ4: Tooling for Control Assurance and Governance	11
2.7 RQ5: Evaluating Auditability and Governance	16
2.8 Synthesis and Related Works	17
3 Methodology	19
3.1 Design Science Context	19
3.2 Regulatory Context and Requirement Derivation	20
3.3 Requirements	20
3.4 Evaluation Framework	21

4	Baseline Pipeline	24
4.1	Overview	24
4.2	Architecture	24
4.3	Functionality	26
4.4	Observed Gaps	27
4.5	Summary	28
5	Proposed Architecture	29
5.1	Lifecycle Overview	29
5.2	Structural Decomposition	31
5.3	Governed Process Maps	34
5.3.1	Change Management Workflow	34
5.3.2	CI/CD Release Workflow	34
5.3.3	Infrastructure Change Workflow	35
5.3.4	Runtime ETL Execution Workflow	36
5.3.5	Lineage and Audit Reconstruction Workflow	37
5.4	Data and Metadata Flows	37
5.4.1	System Data Flow	37
5.4.2	Runtime Data Flow	39
5.4.3	Metadata and Evidence Flow	41
5.5	Worked Example: One Governed EDGAR Run	43
5.6	Summary	45
6	Prototype Implementation	46
7	Evaluation	47
8	Discussion	48
9	Conclusion	49
A	Abbreviations	50
	Bibliography	51

List of Figures

1.1	Gantt chart of project timeline.	3
4.1	Baseline runtime architecture showing orchestration, compute, and storage components.	25
4.2	CI/CD pipeline implemented in GitHub Actions.	26
4.3	Baseline Airflow DAG showing orchestration sequence.	26
4.4	Data path from ingestion to curated output in the baseline pipeline.	27
5.1	Unified lifecycle of the governed ETL architecture.	30
5.2	High-level solution architecture for the proposed pipeline.	32
5.3	Governed change management workflow.	34
5.4	CI/CD workflow for packaging, attestation, and deployment.	35
5.5	Infrastructure change workflow with Terraform state capture.	35
5.6	Governed runtime execution workflow.	36
5.7	Audit reconstruction workflow from dataset to change intent.	37
5.8	System-level data flow across delivery, runtime, and metadata layers.	38
5.9	Runtime flow showing the approved execution context and control gates.	40
5.10	Metadata and evidence flow supporting audit reconstruction.	42

List of Tables

2.1	IaC/PaC tools and the evidence they produce	12
2.2	CI/CD systems as provenance engines	12
2.3	Security scanning tools and audit artefacts	13
2.4	Testing frameworks as reproducible evidence	13
2.5	Container and runtime controls as integrity evidence	13
2.6	Observability stacks and audit support	14
2.7	Orchestration systems and their execution evidence	14
2.8	Processing engines and reproducibility guarantees	15
2.9	Metadata and lineage systems as audit structure	15
2.10	ETL observability as an audit surface	15
3.1	Mapping between toolchain requirements and auditability dimensions.	21
4.1	Representative schema of ingested EDGAR filing metadata.	27

Chapter 1

Introduction

1.1 Context and Motivation

In recent years, the value of data has surpassed that of the software systems built to process it [27, 37]. Organisations increasingly treat data as a product, complete with ownership structures, stewardship responsibilities, and lifecycle management frameworks [27]. This shift has coincided with growing regulatory scrutiny across jurisdictions [27]. Frameworks such as APRA’s Prudential Standards [5–8], alongside global instruments such as the EU’s GDPR and the United States’ HIPAA, demand not only the secure handling of information assets but also transparent processes that demonstrate accountability and control [37].

Extract–Transform–Load (ETL) workflows lie at the heart of this transformation, serving as the operational backbone of analytical systems [36, 38]. As Mahmud and Ikbāl [22] note, modern ETL pipelines have evolved into socio-technical infrastructures that sustain the scalability, reliability, and legitimacy of data-driven decision-making. Yet despite the rise of DataOps, many pipelines remain opaque, lacking end-to-end lineage, reproducibility, and verifiable controls [37]. In regulated environments, such opacity undermines compliance and erodes institutional trust.

Software engineering has already faced similar challenges. DevSecOps — the integration of development, security, and operations — has matured into a discipline embedding verification, automation, and governance across the software delivery lifecycle [3, 4, 10, 28]. If data is now the product, DevSecOps offers an architectural and cultural foundation for treating data pipelines with the same rigour once reserved for software systems [34]. This research explores that proposition: how DevSecOps principles can be adapted to strengthen auditability and data governance within ETL workflows.

1.2 Problem Statement

While the DataOps movement has improved automation and observability, its primary focus remains operational efficiency rather than regulatory assurance [22,25]. Existing orchestration frameworks, though powerful, rarely produce immutable audit trails, reproducible execution states, or cryptographically verifiable lineage. Auditability thus emerges as an accidental property rather than a designed one [39,40]. As a result, organisations face a persistent gap between the technical execution of pipelines and the evidentiary requirements of regulators [6,8].

This thesis addresses that gap. It investigates how a *DevSecOps-enabled toolchain* can enhance auditability and data governance across the ETL lifecycle. DevSecOps has already institutionalised security and compliance for software artifacts, but an equivalent framework for data artifacts has yet to emerge. This absence exposes regulated organisations to compliance risk, weakens governance, and impairs their ability to demonstrate trustworthy data management.

1.3 Research Objectives

This research adopts a Design Science methodology [17] to design, implement, and evaluate a DevSecOps-based toolchain that enhances auditability and governance in ETL workflows. The study follows an iterative cycle of artifact design, prototype construction, and evaluation against explicit metrics for lineage completeness, reproducibility, and tamper detection.

Primary Research Question

How can a DevSecOps toolchain be designed to support ETL workflows in regulated industries, enhancing auditability and data governance?

In the process of answering the above question, this paper explores five sub-questions:

Sub-Questions

1. How do ETL workflows manifest in regulated industries, and what auditability and data governance challenges do they face?
2. How do DevSecOps practices enhance auditability and data governance within the software development lifecycle?
3. In what ways do data lifecycle management practices extend or reshape the software development lifecycle when DevSecOps principles are applied to ETL workflows in regulated environments?

4. What categories of tooling exist in DevSecOps and ETL pipelines, and how do their features support auditability and data governance?
5. How can the effectiveness of a DevSecOps-enabled ETL toolchain be evaluated in terms of auditability and data governance?

1.4 Contributions

This thesis makes three primary contributions. First, it proposes an **architectural model** for a DevSecOps-enabled ETL toolchain, aligning the data lifecycle with secure software delivery practices to improve auditability and governance. Second, it presents a **prototype implementation** that integrates infrastructure as code, automated security scanning, and immutable logging into an operational ETL environment built with Airflow, Spark, and AWS components. Third, it defines an **evaluation framework** for characterising auditability and governance through traceability, reproducibility, and tamper-resistance criteria, directly addressing Sub-Question 5. Together, these contributions demonstrate how DevSecOps principles can be embedded in data-centric systems to enhance compliance and trust.

1.5 Project Timeline




Figure 1.1: Gantt chart of project timeline.

1.6 Thesis Structure

The remainder of this thesis is organised as follows:

- **Chapter 2** reviews literature on ETL pipelines, data governance, and DevSecOps practices.

- **Chapter 3** explains the Design Science methodology, design constraints, and evaluation criteria.
- **Chapter 4** describes the baseline ETL pipeline and identifies auditability and governance gaps.
- **Chapter 5** presents the proposed DevSecOps-integrated architecture.
- **Chapter 6** outlines the prototype implementation and toolchain configuration.
- **Chapter 7** evaluates the system against the auditability and governance criteria defined in Chapter 3, addressing Sub-Question 5.
- **Chapter 8** discusses implications, limitations, and avenues for future research.
- **Chapter 9** concludes with reflections on how DevSecOps can enable trustworthy data governance in regulated environments.

The Preliminary results for this Thesis are split between the evaluation framework found in Chapter 3 and the base ETL pipeline described in Chapter 4.

Chapter 2

Literature Review

This chapter reviews the body of knowledge relevant to the design of ETL pipelines in regulated industries and the application of DevSecOps principles. It begins by outlining the context of ETL and governance, defines the key concepts central to this research, and then examines the literature associated with each research question.

2.1 Background

Data pipelines underpin modern analytical systems, forming the connective infrastructure between data sources and analytic platforms [34]. They transform raw data into governed, traceable data products used in applications such as fraud detection, market analysis, compliance reporting, and metadata processing.

Although the extract–transform–load (ETL) process appears simple in theory, it is in practice complex, resource-intensive, and often the most time-consuming element of data warehousing [18]. ETL pipelines act as the “eyes of the system,” feeding reliable data downstream; when lineage is lost or visibility fails, the reliability of the entire process is compromised [36]. Vayyala [36] emphasises that maintaining end-to-end lineage across transformations is vital for sustaining trust in analytics-driven decisions.

In regulated industries, compliance depends as much on how data is processed as on the outcomes produced. Frameworks such as GDPR, SOX, and APRA’s CPS standards require secure, ethical, and well-documented data handling [5–8, 27]. These frameworks place accountability, ownership, and an end-to-end evidence chain at the centre of compliance culture.

Auditability, in this context, is the ability to demonstrate compliance through reproducibility, lineage, and tamper-evident records [37]. Trustworthy analytics rely not only on integrity of data but also on assurance evidence of correct process execution. Governance provides the organisational scaffolding for this assurance—linking policy to technical enforcement through defined roles, procedures, and controls [27, 37]. Failures in either dimension expose organisations to legal,

financial, and reputational risk [27].

DevOps emerged to accelerate software delivery through automation, continuous integration and deployment (CI/CD), and rapid feedback cycles [28]. DevSecOps extends this model by embedding security and compliance across the software development lifecycle. Practices such as automated testing, compliance-as-code, artifact signing, and tamper-evident logging produce assurance evidence that can be inspected by regulators [4, 10, 28]. This convergence bridges engineering and governance concerns, establishing control assurance as a built-in property of the delivery process.

As data itself becomes a primary strategic asset, surpassing the value of the systems that process it [27, 37], the mechanisms of DevSecOps must be extended to the data pipelines that sustain organisational decision-making.

2.2 Defining Key Concepts

The following subsections define the key concepts underpinning this research: DevSecOps, ETL, data governance, and auditability. Each definition draws from existing literature and sets the theoretical foundation for later analysis.

DevSecOps

DevSecOps is the proactive and continuous enforcement of security and compliance requirements throughout the software development lifecycle (SDLC) [4]. Unlike traditional models that treat security as an external stage or afterthought, DevSecOps embeds it within every phase of the lifecycle [3]. Compliance becomes the expected outcome by default, and deviations from defined standards must be explicitly justified [10].

DevSecOps is notoriously difficult to define succinctly. It is often described as the sum of its implementation patterns—CI/CD, infrastructure-as-code (IaC), policy-as-code (PaC), software scanning, automated tests, and GitOps. A more practical lens is to examine the gap created when DevSecOps is removed: the un-governed space between code that compiles and a production service. In that gap, identity and approval are ad hoc, policy enforcement is manual, change control is episodic, evidence is incomplete, and configuration drifts without reconciliation. DevSecOps closes this gap by making control assurance continuous, placing verification at the point of change, and generating the lineage and evidence metadata needed for independent assessment [4, 33].

A defining characteristic of DevSecOps is its reliance on automation to enforce non-functional requirements such as control assurance and governance. Through automated testing, provisioning, monitoring, and policy enforcement, systems produce assurance evidence—signed artifacts, tamper-evident records, and control indicators—that regulators and auditors can independently assess [4, 44]. This

automated assurance narrows the gap between delivery teams and governance stakeholders [3].

By combining cultural expectations of shared responsibility with technical mechanisms for verification, DevSecOps establishes a delivery model in which security and compliance are continuous, measurable, and intrinsic—rather than externally imposed.

ETL

Extract–Transform–Load (ETL) is a structured process that ingests, transforms, and persists data into target repositories such as data warehouses, lakes, or file stores [18]. It traditionally comprises three stages: extraction, where data is sourced from heterogeneous systems; transformation, where it is cleansed, standardised, integrated, and reshaped into coherent semantic structures; and loading, where transformed data is stored for downstream analysis [18]. Kimball and Caserta describe ETL as both the most resource-intensive and most critical link between operational data and analytical use.

Modern practice extends beyond this classical batch model. Approaches such as extract–load–transform (ELT) and hybrid streaming–batch architectures leverage distributed and cloud-native infrastructure [22]. ETL now encompasses not only data integration and quality but also governance and compliance. In regulated industries, pipelines must guarantee lineage, reproducibility, and control assurance in addition to performance and scalability [22].

ETL is thus best understood as a flexible socio-technical process. It concludes when data is made persistent and reusable but intersects with broader lifecycle concerns such as orchestration, monitoring, and policy enforcement. In this sense, ETL pipelines underpin both analytical scalability and institutional trust [22].

Data Governance

Data governance provides the strategic framework that defines how data is collected, managed, and protected across an organisation. It establishes the policies, roles, standards, and processes that ensure information assets remain accurate, consistent, secure, and compliant with internal and external regulations [27, 37]. Nai, Rifai, and Sadiq [27] describe governance as a structured framework that ensures the availability, integrity, usability, and security of data through stewardship and accountability roles.

Vayyala [37] expands this understanding by positioning governance as the foundation of reliable data ecosystems—linking policy to technical enforcement through automation, lineage tracking, and tamper-evident records. In practice, governance institutionalises accountability so data products can withstand regulatory scrutiny.

Auditability

Auditability is the designed capability of a system to allow its processes, controls, and outcomes to be observed and verified by stakeholders. Weigand and Elsas [40] define auditability through five interdependent requirements:

- **Normativity:** a clear regulatory or contractual framework that governs operations.
- **Accountability:** verifiable statements linking actions to responsible agents and governing norms.
- **Referentiality Infrastructure:** records that connect operations to their original points of recording.
- **Reliability:** faithful and tamper-resistant representation of events.
- **Verifiability:** the ability of external parties to independently test or validate claims.

Although these principles originate in organisational and financial control, they can be transferred to information systems. Extending this theoretical foundation, Wang et al. [39] demonstrate how auditability can be operationalised through third-party verification, cryptographic proofs of integrity, and support for dynamic data operations. Their work illustrates that auditability is not merely an organisational goal but a property that can be engineered into data management infrastructures.

Applied to ETL pipelines and DevSecOps practices, these five dimensions map directly to regulated-industry concerns:

- **Normativity:** embedding compliance and data-protection policies within workflow execution.
- **Accountability:** linking each transformation or deployment to authenticated agents or automated processes.
- **Referentiality Infrastructure:** ensuring full lineage tracking across extraction, transformation, and loading.
- **Reliability:** maintaining tamper-evident records, signed artifacts, and cryptographic hashes.
- **Verifiability:** enabling reproducibility tests and independent control testing via CI/CD-integrated checks.

Together, these adapted requirements frame auditability as a designed and engineered property of ETL systems. Every data operation and system change must attach to an end-to-end evidence chain that is reproducible and independently verifiable against regulatory expectations.

2.3 RQ1: ETL Pipelines in Regulated Industries

ETL pipelines in regulated industries operate within a constant tension between agility and accountability. They are expected to deliver analytical value with speed and scale, yet must simultaneously produce verifiable evidence that every data movement complies with institutional and regulatory constraints. The modern pipeline has therefore evolved into a socio-technical infrastructure—one that mediates trust between engineers, auditors, and regulators [22,36]. Its success cannot be judged solely by throughput or latency; rather, by its ability to maintain lineage, transparency, and compliance under continuous change.

Financial, healthcare, and critical-infrastructure systems exemplify this dual demand. Frameworks such as APRA’s CPS 230, CPS 234, and CPS 510 [6–8] institutionalise separation of duties and controlled promotion of change, translating ethical responsibility into technical procedure. These standards do not exist apart from engineering practice—they shape it—embedding legal accountability into operational design. Governance scholarship reaches similar conclusions, emphasising that tamper-evident logging, role-based accountability, and retention aligned with compliance lifecycles form the practical machinery of institutional trust [27, 37, 38]. Cloud-native orchestrators such as Airflow, AWS Glue, and Azure Data Factory provide elasticity and automation, but they operate under the shadow of these mandates. Their convenience is constrained by the need to preserve traceability and access control consistent with CPG 235 [5]. In this context, orchestration itself becomes a form of governance: a choreography of tasks and dependencies that leaves a verifiable trail of execution [31].

Over time, this operational discipline has turned the pipeline from a means of data transport into a living archive of control assurance. Continuous validation, versioning, and policy enforcement integrate assurance logic directly into the runtime environment [37]. Metadata management forms the connective tissue between policy and process, ensuring that lineage and evidence metadata, schema evolution, and transformation logic can be reconstructed and defended [38]. Yet even within mature systems, full control assurance remains elusive. Toolchain heterogeneity fragments lineage across layers of orchestration, computation, and storage. Business logic embedded in SQL or user-defined functions escapes formal tracing, while schema drift and inconsistent metadata propagation corrode reliability [22]. Few frameworks create tamper-evident, cryptographically verifiable artifacts by default [38,40]. What is missing is not information, but integrity—the ability to prove that recorded operations are both complete and untampered. The fragility of this end-to-end evidence chain provides the practical motivation for integrating DevSecOps assurance into ETL environments, turning validation from an external audit to an intrinsic system property.

2.4 RQ2: DevSecOps and Governance in the SDLC

The software development lifecycle (SDLC) treats change as a controlled and observable process. Standards such as ISO/IEC/IEEE 12207 and 15288 define the artifacts, procedures, and documentation that make each phase accountable [1, 16]. Governance in this model is not an external overlay—it is embedded in the way systems are planned, built, and released.

DevSecOps strengthens this principle by embedding assurance, security, and compliance across every phase of delivery [4, 33]. The delivery pipeline becomes a generator of assurance evidence: each build, test, and deployment produces signed artifacts, tamper-evident records, and compliance manifests that bind design intent to execution [3, 4, 11, 40]. These artifacts collectively form a continuous end-to-end evidence chain that makes system state traceable and verifiable at any point in time.

Automation enables this continuity. Infrastructure-as-code (IaC), continuous-integration/continuous-delivery (CI/CD), and policy-as-code (PaC) apply verification directly at the point of change. Static and dynamic testing, dependency scanning, and approval gates transform compliance into an embedded control validation signal rather than a periodic audit [3, 16, 21]. The NIST Secure Software Development Framework codifies this expectation: traceable metadata, lineage and evidence metadata, and policy enforcement are standard lifecycle outputs [33]. Mature build systems act as lineage and evidence-metadata engines, linking requirements, commits, and runtime behaviour within the same evidentiary record.

As software supply chains expand, continuous verification of component origin and integrity becomes essential [21, 33]. Toolchains respond through automated scanning, provenance analysis, and AI-assisted policy enforcement [2, 23]. Requirement traceability then links regulatory expectations directly to runtime artifacts [1, 11], ensuring that what is built, deployed, and reviewed aligns to the same authoritative record [28].

Several implementation patterns operationalise DevSecOps principles of automation and assurance. Infrastructure-as-Code (IaC) enforces reproducible, governed environments; Policy-as-Code (PaC) embeds compliance validation within delivery pipelines; and GitOps extends these mechanisms into runtime operations by reconciling deployed state with version-controlled configuration. Together, these patterns close the loop between design-time verification and operational control, forming the practical foundation of continuous control monitoring within the SDLC.

2.5 RQ3: Comparing the Data Lifecycle and the SDLC

The data lifecycle (DLC) aims for control and reliability but lacks the SDLC's formal standardisation. Wing describes the DLC as a dynamic, context-dependent model rather than a fixed sequence of phases [41]. Most frameworks share the same goal: to keep data reliable, accessible, and verifiable throughout its lifespan [9,20]. The SDLC governs change and transformation; the DLC governs persistence and stewardship. One manages evolution, the other preserves continuity.

Despite their structural differences, the two lifecycles can be aligned. The Data Management Body of Knowledge (DAMA-DMBOK) provides one of the most widely adopted codifications of the DLC, framing governance through roles, metadata catalogues, lineage records, and quality assessments that together establish an end-to-end evidence chain [13]. These governance artifacts closely mirror the information items defined in ISO/IEC/IEEE 15289 [1], demonstrating that the logic of software assurance maps naturally onto data management practice. DevSecOps supplies the mechanisms for this connection: continuous validation and automated security controls apply equally to code and data [4,33]. Lineage manifests, schema definitions, and transformation scripts can therefore be versioned and audited alongside software components.

Encryption, secrets management, and role-based access add cryptographic assurance to this framework [29,34]. Agile governance pushes data validation earlier in the lifecycle [19,36,38], mirroring the shift in DevSecOps toward early verification. These expectations reflect APRA's prudential standards [5–8], where assurance is evidenced through the operational trace rather than through static reports.

When the SDLC and DLC intersect, they create a shared evidentiary framework. ISO/IEC/IEEE 15289's artifacts—plans, specifications, procedures, and reports—extend to include lineage documentation, schema evolution, and validation outputs. IaC and policy-as-code make these records executable [3,26,28]. Reproducibility, tamper detection, and governance then operate through the same automated pipelines that deliver software.

In regulated ETL workflows, control assurance becomes a first-class design property. Each transformation of code or data leaves a signed, tamper-evident record that shows what occurred, who performed it, and why it was authorised.

2.6 RQ4: Tooling for Control Assurance and Governance

Assurance depends on implementation. The ability to trace, verify, and reproduce each operational event relies on tools that enforce control, capture lineage, and preserve evidence of change. In both DevSecOps and ETL environments, tooling

translates regulatory requirements such as those in CPG 235 and CPS 234/510 into machine-verifiable artefacts [5–7].

This section examines the tool categories that make such evidence possible. In DevSecOps, frameworks for infrastructure-as-code, policy validation, and continuous integration generate signed artefacts and lineage and evidence metadata that prove system integrity. In ETL pipelines, ingestion and orchestration components, processing engines, metadata catalogues, validation frameworks, and monitoring systems create traceable datasets and transformation histories that support data governance. Each tool category contributes to the same end-to-end evidence chain, linking software and data processes into a single, verifiable system of record.

DevSecOps tooling: from code to operations

Infrastructure- and Policy-as-Code

IaC and PaC convert infrastructure and policy into versioned artefacts so that compliance checks run before change. Terraform, CloudFormation, OPA, Checkov, and Conftest generate explicit evidence of desired state and policy evaluation [4, 29]. Key contrasts appear in Table 2.1.

Table 2.1: IaC/PaC tools and the evidence they produce

Tool	Strengths	Limitations
Terraform	Versioned infra; plan/apply logs; OPA integration	Limited audit trail
CloudFormation	CloudTrail linkage; tight IAM controls	Vendor lock-in
OPA / Conftest	Declarative policy; decision logs as evidence	Rule complexity

CI/CD orchestration

GitHub Actions, GitLab CI, Jenkins, and Azure Pipelines attach identity, review, and signatures to every run; pipeline logs form part of the evidentiary chain [11, 33]. Table 2.2 summarises key trade-offs.

Table 2.2: CI/CD systems as provenance engines

Tool	Strengths	Limitations
GitHub Actions	Signed runs and artefacts; Git provenance	Coarse access audit
GitLab CI	Built-in approvals; compliance reports	Higher operational cost
Jenkins	Extensible audit plugins	Plugin sprawl
Azure Pipelines	RBAC tie-in; policy checks	Cloud-specific integrations

SAST, DAST, and SCA

Security scanners create structured evidence of vulnerabilities and dependency lineage. They integrate with CI/CD pipelines and feed policy gates [12, 29, 32]. Table 2.3 summarises trade-offs.

Table 2.3: Security scanning tools and audit artefacts

Tool	Strengths	Limitations
SonarQube	Tracks code quality and issues	Limited container context
OWASP ZAP	Reproducible web tests; open output	Slow at scale
Snyk	Dependency risk scoring; SBOMs	Proprietary backend
Trivy	Combined SCA/SAST; SBOMs	Narrow policy scope

Testing and validation

Test frameworks provide repeatable proof of expected behaviour; their outputs become verifiable artefacts that link requirements to executions [32, 33]. Table 2.4 outlines key examples.

Table 2.4: Testing frameworks as reproducible evidence

Framework	Strengths	Limitations
Postman / Newman	Versioned API suites; runnable logs	Manual upkeep of collections
PyTest	Parametrised fixtures; stable outputs	Python-only scope
JUnit	Mature CI integration; clear reports	No native audit metadata

Containers and runtime

Images, admission controls, and runtime events contribute directly to system integrity evidence [4, 28]. Table 2.5 lists key tools and their audit characteristics.

Table 2.5: Container and runtime controls as integrity evidence

Tool	Strengths	Limitations
Docker	Image signatures; digest verification	Limited runtime policy control
Kubernetes	Declarative policy; drift detection	Complex multi-cluster audits
Falco	Real-time event monitoring	High rule upkeep
Prisma Cloud / Sysdig	Centralised compliance view	Proprietary; costly

Monitoring and observability

Control indicators and logs form the final evidence layer, supporting incident reconstruction and trend analysis [4, 29]. Table 2.6 highlights typical stacks.

Table 2.6: Observability stacks and audit support

Tool	Strengths	Limitations
Prometheus	Transparent, timestamped metrics	High cost for long retention
Grafana	Clear dashboards; evidence views	Dependent on source data quality
ELK / Splunk	Searchable forensic logs	Operational cost; index complexity

ETL tooling: from ingestion to lineage

ETL ecosystems expose their own audit surface. Ingestion tools capture how data enters governed infrastructure; orchestrators record how it moves; processing engines determine how it changes; and metadata, validation, privacy, and monitoring systems preserve the story of what happened and when. Together, these tools produce the evidence needed to link data handling to governance expectations in regulated environments [22, 35, 38, 42].

Orchestration

Orchestrators manage dependencies, retries, and execution order. Their run histories and task states give ETL workflows a temporal structure that can be replayed and inspected [22, 35]. They connect ingestion with transformation and expose where control decisions were applied. Table 2.7 compares common orchestration systems and the forms of execution evidence they provide.

Table 2.7: Orchestration systems and their execution evidence

Tool	Strengths	Limitations
Apache Airflow	Clear DAGs; task history	Manual scaling; weak lineage
Azure Data Factory	Built-in control; Purview link	Limited cross-cloud use
AWS Step Functions	Visual flows; CloudWatch logs	Coarse logs; AWS-only
Dagster	Data-aware; built-in lineage	Small community; steep setup

Transformation engines

Processing engines carry out the transformations that regulators most care about; they decide how raw records become derived data products. Engines such as Spark, EMR, Dataflow, and Delta Lake can provide deterministic checkpoints

and time-travel capabilities that support reproducible proofs of change [22, 24]. Table 2.8 outlines their main assurance properties.

Table 2.8: Processing engines and reproducibility guarantees

Engine	Strengths	Limitations
Apache Spark	Distributed compute; checkpoint versioning	Complex lineage tracing
Google Dataflow	Managed scaling; audit logging	Cloud lock-in
Delta Lake	ACID tables; rollback support	Storage overhead

Metadata and lineage

Catalogues and lineage systems form the spine of auditability: they connect datasets, schemas, and processes into a coherent graph that can be queried and tested [38, 42]. These tools turn implicit relationships into explicit evidence, supporting both governance and impact analysis. Table 2.9 summarises key options.

Table 2.9: Metadata and lineage systems as audit structure

Tool	Strengths	Limitations
Apache Atlas	Open lineage; policy integration	Latency at scale
Microsoft Purview	Rich lineage views; Azure native	Proprietary scope
Amundsen	Lightweight metadata discovery	Limited enforcement links

Operational monitoring

Monitoring and observability tools turn runtime behaviour into an audit surface. Metrics, traces, and logs enable lineage verification, incident reconstruction, and control testing over time [35]. In ETL contexts, these systems complete the evidence chain by linking pipeline events to infrastructure and application signals. Table 2.10 summarises common options.

Table 2.10: ETL observability as an audit surface

Tool	Strengths	Limitations
Azure Monitor	Unified infra and data view	Short retention limits
AWS CloudWatch	Deep AWS integration; trace logging	Complex audit queries
Prometheus	Open telemetry; clear metrics	High data volume

Convergence

GitOps as an Operational Assurance Pattern

GitOps is a practical implementation of DevSecOps principles that operationalises configuration and change assurance. By declaring desired system state in version control and automating reconciliation, GitOps enforces configuration provenance, segregation of duties, and continuous control validation across environments. In this unified model, the same controls that attest software integrity also attest data integrity. Approval gates in CI/CD govern the promotion of ETL workflows; policy checks constrain data transfers; and container attestations record the run-time context of both applications and transformations. Observability systems link these artifacts into a continuous audit timeline spanning code, infrastructure, and data [4, 15, 33]. GitOps enforces configuration provenance; complementary lineage controls extend this assurance to data artifacts and transformations.

These implementation patterns—particularly GitOps—form the operational basis of the proposed architecture in Chapter 5, where declarative configuration and automated reconciliation underpin continuous assurance across both software and data lifecycles. Regulatory expectations under CPG 235 and CPS 234/510 [6, 7] are satisfied not through separate audits but through a single, declarative lifecycle.

Three assurance patterns make this convergence practical. *Single-path promotion* ensures that every artifact—application, configuration, or dataset—follows the same reviewed path to production. *Policy-guarded reconciliation* binds control objectives to each applied change. *Attested runtime* couples container and orchestration evidence to specific commits and policies, closing the loop between declared intent and observed outcome [4, 33].

2.7 RQ5: Evaluating Auditability and Governance

Evaluation forms a core component of Design Science Research, providing the means to demonstrate an artifact’s utility, quality, and validity in context. Hevner [17] and Peffers [30] describe evaluation as integral to the build–evaluate cycle, while Gregor and Hevner [14] distinguish between *ex-ante* evaluation—assessing design logic before implementation—and *ex-post* evaluation, which measures performance in operation. Together, these works establish evaluation as both a validation activity and a source of knowledge in its own right.

Within DevSecOps literature, evaluation typically focuses on security or compliance performance rather than evidentiary assurance. Souppaya and Scarfone’s NIST SSDF [33] and Anisetti et al. [4] present process metrics for secure software delivery, while Weigand and Elsas [40] frame auditability as a design property requiring verifiable, tamper-evident records. These studies demonstrate that control assurance can be measured, but they stop short of applying such evaluation to

data-intensive workflows that must also satisfy regulatory expectations for lineage, reproducibility, and provenance.

From a data-engineering perspective, evaluation of auditability is most often addressed indirectly through lineage reconstruction, data quality assessment, and provenance capture [22, 36, 38]. ETL research typically measures performance in terms of latency, scalability, or transformation accuracy, while evidence of compliance and reproducibility is treated as secondary. Existing frameworks for data lineage and metadata management establish mechanisms for tracing data flow but rarely define how the completeness or trustworthiness of that lineage should be evaluated [39]. As a result, few approaches combine the quantitative rigor of DevSecOps metrics with the contextual assurance required in regulated data environments.

The literature therefore points to a gap between existing DevSecOps evaluation frameworks—focused on software artifacts—and the needs of data-governance environments, where evidentiary completeness and verifiable lineage are central. Addressing this gap requires extending evaluation beyond security metrics to encompass measurable dimensions of auditability and governance in data-centric systems.

2.8 Synthesis and Related Works

Research at the intersection of DevSecOps and data governance converges on the same foundation—automation, traceability, and verifiable control—but few works treat these as components of a unified evidence system spanning the full ETL lifecycle. Existing frameworks address individual facets of the assurance problem—data quality, lineage, or security—but seldom integrate them under a common governance architecture capable of satisfying regulatory expectations for auditability.

The IGI Global collection *Data Governance, DevSecOps, and Advancements in Modern Software* [27] provides the strongest conceptual foundation. Its chapters illustrate how metadata management, lineage, and policy-as-code can bridge security automation and governance accountability. Vayyala [38] presents metadata as a governance mechanism; Kose [19] links agile compliance to DevSecOps; and Soussane [34] applies these principles to real-time analytics. Collectively, they demonstrate how DevSecOps can embed governance within delivery pipelines, though they stop short of implementing tamper-evident and reproducible evidence across the data lifecycle.

Subsequent studies in software assurance and lifecycle management reinforce these ideas but remain domain-bound—software or data, not both. Anisetti et al. [4] and Souppaya and Scarfone [33] formalise metrics for secure software delivery, while Wing [41] and Zhang et al. [43] explore stewardship and reliability within data lifecycles. Yet no current model provides a framework that spans both domains to create a single, tamper-evident audit trail linking code, configuration,

and data artifacts.

The research gap is therefore twofold. First, auditability remains treated as an emergent property rather than a designed capability of ETL systems. Second, evaluation methods in both DevSecOps and data governance literature rarely consider evidentiary completeness or lineage trustworthiness as measurable outcomes. This thesis extends the existing body of work by proposing a unified DevSecOps-enabled ETL toolchain that operationalises continuous assurance through verifiable lineage, reproducibility, and governance evaluation.

Chapter 3

Methodology

This chapter outlines the methodology used to design, construct, and evaluate a DevSecOps toolchain that governs ETL workflows in regulated environments, including the framework for evaluating auditability and governance defined in Sub-Question 5.

The research applies a Design Science Research (DSR) approach, chosen for its focus on building artifacts that address real-world problems while generating transferable knowledge [17]. Rather than developing a new ETL pipeline, this study designs the system that governs such pipelines—the architecture, control automation, and controls that embed assurance and compliance into operational data workflows.

3.1 Design Science Context

Design Science generates knowledge through the purposeful creation and evaluation of artifacts [30]. It follows an iterative build–evaluate process in which understanding arises from constructing, implementing, and refining solutions. In this study, the artifact is a *DevSecOps governance and assurance framework*—a set of integrated components that enforce policy, record lineage and evidence metadata, and assure integrity across an ETL pipeline. The pipeline itself functions as a *reference pipeline* that represents a typical control-baseline environment within a regulated domain.

This methodological choice is appropriate because Design Science balances *rigour*—alignment with regulatory and theoretical principles—with *relevance*—demonstrated utility in practice. Each design iteration: problem identification, artifact construction, demonstration, and evaluation, corresponds to a development stage where the framework’s assurance features are incrementally refined and tested. This cyclical process mirrors the continuous control monitoring model advocated in DevSecOps, where systems evolve through assurance evidence and feedback.

3.2 Regulatory Context and Requirement Derivation

The toolchain design is guided by the Australian Prudential Regulation Authority's (APRA) Prudential Standards: CPS 230 (*Operational Risk Management*), CPS 234 (*Information Security*), and CPS 510 (*Governance*) [6–8]. These standards define what regulated entities must achieve in governance, security, and resilience, but they do not prescribe how those objectives should be implemented technically. Their descriptive, outcomes-based language demands interpretation into enforceable engineering controls.

Although the methodology is domain-agnostic, this research focuses on the *financial sector*. Finance offers a mature and well-documented regulatory landscape that provides both conceptual depth and practical relevance for governance evaluation. APRA's prudential standards form a coherent, enforceable set of requirements that align with the objectives of assurance, lineage, and accountability identified in the literature. This focus reflects the local regulatory environment and the availability of structured compliance expectations that make the sector suitable for Design Science evaluation.

Accordingly, the scope is narrowed to a practical subset of CPS requirements most relevant to data-pipeline governance. The selected principles—continuity, accountability, and control—are translated into implementable requirements for control automation and evidence generation. This refinement is informed by the literature review in Chapter 2, which frames auditability (as defined there) in terms of lineage, evidence, and verification across socio-technical systems.

3.3 Requirements

The following requirements define what the DevSecOps framework must achieve to operationalise prudential expectations. They establish the conditions under which the framework can govern a pipeline by embedding control, lineage, and continuous assurance.

1. **R1: Governance Accountability (CPS 510).** The framework must preserve authorship and approval lineage for all operational and configuration changes, ensuring full traceability of decision authority and execution.
2. **R2: Information Security Assurance (CPS 234).** The framework must maintain artifact integrity and confidentiality through tamper-evident storage, cryptographic verification, and continuous security scanning (SCA, SAST, and DAST).
3. **R3: Operational Risk Control (CPS 230).** The framework must detect, record, and remediate process or configuration failures, maintaining reproducibility and recovery within defined assurance thresholds.

4. **R4: Continuous Control Monitoring (CPS 234, 510).** The framework must automate policy validation and control testing in the CI/CD process, producing machine-verifiable assurance evidence through policy-as-code enforcement.
5. **R5: End-to-End Evidence Chain (CPS 230).** The framework must generate and retain a tamper-evident, queryable record that captures the lineage of code, configuration, and data artifacts throughout the system lifecycle.

These requirements shift focus from the business logic of ETL processes to the *meta-layer* of assurance—how evidence is produced, tested, and preserved. They align with the five auditability dimensions defined in Section 2.2: normativity, accountability, referentiality, reliability, and verifiability.

Table 3.1: Mapping between toolchain requirements and auditability dimensions.

Requirement	Normativity	Accountability	Referentiality	Reliability	Verifiability
R1 Governance Accountability	●	●			
R2 Information Security Assurance				●	●
R3 Operational Risk Control			●	●	
R4 Continuous Control Monitoring	●	●			
R5 End-to-End Evidence Chain			●		●

Governance and control monitoring (R1, R4) express the normative and accountable aspects of control; security and risk management (R2, R3) reinforce reliability and traceability; and the end-to-end evidence chain (R5) provides verifiable, referential links across artifacts and executions. Together, these create a governance and assurance framework where assurance is continuous and evidence is intrinsic.

3.4 Evaluation Framework

Evaluation in Design Science verifies that an artifact fulfils its intended purpose. In this study, evaluation operationalises Sub-Question 5 by defining how the DevSecOps framework will be assessed against the requirements R1–R5. At this preliminary stage, evaluation criteria are expressed qualitatively rather than as fixed numeric thresholds. Each requirement becomes an evaluation unit with observable control indicators, data sources, and assurance conditions, while quantitative benchmarks are deferred to Thesis B once empirical measurements are available.

R1 Governance Accountability

Objective: Confirm that all changes are traceable to authenticated identities and documented approvals.

Method: Examine change events, approval records, and version metadata for completeness and signature validity.

Control Indicators:

- Configuration and code changes are traceable to verified identities.
- Approval events are associated with tamper-evident artifacts.
- Lineage between versions remains continuous and reconstructable.

Auditability Dimensions: Normativity, Accountability.

R2 Information Security Assurance

Objective: Demonstrate that artifacts and configurations are verified for integrity and security.

Method: Conduct automated and manual reviews confirming that artifacts are signed, validated, and free from known vulnerabilities.

Control Indicators:

- Artifacts are verified for integrity through signatures or checksums before use.
- No integrity violations remain undetected or unremediated during testing.
- Verification records are retained and accessible for independent review.

Auditability Dimensions: Reliability, Verifiability.

R3 Operational Risk Control

Objective: Validate reproducibility and recovery from prior workflow states.

Method: Execute controlled replays of historical states and compare outputs to reference baselines.

Control Indicators:

- Lineage information is sufficiently complete to reconstruct mappings between inputs, processing steps, and outputs.
- Historical workflow states can be replayed and produce consistent, explainable results.
- Recovery procedures restore prior states within limits that are operationally acceptable for the reference context.

Auditability Dimensions: Referentiality, Reliability.

R4 Continuous Control Monitoring

Objective: Verify that compliance enforcement and validation operate automatically for every change.

Method: Observe multiple workflow executions to confirm policy checks and validations occur consistently.

Control Indicators:

- Policy checks and control validations run consistently for each change and execution path.
- Unverified or unapproved changes are prevented from reaching production environments.
- Each control evaluation produces verifiable evidence of the checks performed and their outcomes.

Auditability Dimensions: Normativity, Accountability.

R5 End-to-End Evidence Chain

Objective: Confirm that the framework maintains a tamper-evident, queryable record of artifacts and lineage.

Method: Perform end-to-end tracing of artifacts across executions to ensure each output can be reconstructed from recorded evidence.

Control Indicators:

- Evidence and lineage records validate successfully across representative workflow executions.
- Each transformation is traceable to its original inputs and associated configurations.
- No unexplained inconsistencies appear in signature or hash verification during evaluation.

Auditability Dimensions: Referentiality, Verifiability.

Evaluation Synthesis

Each requirement is tested through direct observation and analysis, with results mapped back to the auditability dimensions in Table 3.1. Collectively, these evaluations determine whether the governance and assurance framework establishes a verifiable end-to-end evidence chain across its lifecycle, thereby addressing Sub-Question 5 in the context of the reference pipeline.

Chapter 4

Baseline Pipeline

This chapter defines the baseline pipeline that serves as the control model for later evaluation. It represents a minimal but coherent ETL workflow constructed from standard open-source and cloud-native components, configured without embedded governance or assurance controls. The baseline shows how data moves, where control resides, and what traces of evidence are naturally produced. It establishes the reference point against which the DevSecOps architecture in Chapter 5 will be measured.

4.1 Overview

The baseline pipeline ingests public financial filings from the SEC EDGAR API, stages them as raw CSV objects in Amazon S3, transforms them with PySpark on EMR, and outputs curated Parquet datasets back to S3. Airflow, running on a Dockerised EC2 instance, orchestrates the workflow. The infrastructure is declared in Terraform and deployed through GitHub Actions, which performs linting, builds the Airflow image, provisions resources, and redeploys on each `push` to `main`.

Conceptual scope. The configuration models a typical cloud-native pipeline stripped of security and compliance instrumentation. Defaults are accepted to keep the environment simple and reproducible; the goal is not performance optimisation but clarity. This baseline represents the state of practice before DevSecOps principles are introduced to turn operational consistency into demonstrable assurance.

4.2 Architecture

The system operates across three planes—*orchestration*, *compute*, and *storage*—with automation and infrastructure layers supporting them. Figure 4.1 shows the overall arrangement within the AWS environment. Each component reflects a tooling

category defined in Chapter 2 and illustrates both its operational value and its assurance gap.




Figure 4.1: Baseline runtime architecture showing orchestration, compute, and storage components.

Orchestration: Airflow. Airflow defines task dependencies and execution order across the ETL process. It embodies the orchestration tools discussed in Section 2.6, functioning as a runtime scheduler and provenance log. In the baseline, these logs are mutable text records with no signatures or policy linkage; they record execution but not authority. The system therefore guarantees temporal order but not accountable control.

Compute: PySpark on EMR. PySpark provides distributed transformation aligned with the compute and runtime control category from Section 2.6. It enforces schema validation and predictable transformations across nodes. Yet job scripts, dependencies, and configurations remain mutable; no digital signatures or versioned attestations confirm authorship. The compute layer executes correctly but offers no cryptographic proof of integrity.

Storage: Amazon S3. S3 acts as both staging and curated store. It ensures durability and timestamped versioning but offers no immutability or content signing. Objects can be modified or deleted without trace, forcing lineage reconstruction from logs rather than metadata. The storage layer is reliable, yet unverifiable.

Infrastructure: Terraform. Terraform declares the infrastructure state and belongs to the Infrastructure-as-Code and Policy-as-Code tooling class. It provides repeatability through declarative configuration but lacks persistence in its execution records; logs are transient and unsigned. The result is infrastructure that can be reproduced, but not proven to match the reviewed plan.

Automation: GitHub Actions. Continuous integration and delivery are handled through GitHub Actions, shown in Figure 4.2. This pipeline unifies code, configuration, and deployment within a single repository; it represents the first step toward a GitOps model. However, it omits policy validation, multi-party approval, and artifact signing. Each push triggers infrastructure updates automatically, without checks that would confirm authenticity or intent.




Figure 4.2: CI/CD pipeline implemented in GitHub Actions.

4.3 Functionality

Figure 4.3 shows the Airflow DAG that governs workflow execution. Tasks proceed linearly—ingestion, transformation, and manifest recording. The process runs reliably and can be repeated, but it cannot be independently verified.




Figure 4.3: Baseline Airflow DAG showing orchestration sequence.

Ingestion. A Python task retrieves structured JSON filings from the EDGAR API, normalises them, and writes timestamped CSV files to a raw S3 bucket. Each run

creates a new partition, forming the input dataset for downstream processing. The representative schema is shown in Table 4.1.

Field	Type	Description
cik	String	Company identifier assigned by the SEC.
form_type	String	Filing category (e.g., 10-K, 8-K).
filing_date	Date	Official submission date.
accession_no	String	Unique accession number for each filing.
file_url	String	Direct URL to the source document.
company_name	String	Registered name of the filer.

Table 4.1: Representative schema of ingested EDGAR filing metadata.

Transformation. The PySpark job reads the raw objects, validates schema consistency, deduplicates, and writes partitioned Parquet datasets to the curated bucket. Logs capture execution time and row counts but omit lineage metadata linking inputs, scripts, and outputs. The pipeline delivers correct results without proving how they were produced.

Output and manifest. A final task generates a `_SUCCESS` manifest in S3, listing output paths for verification. The manifest records completion but not content integrity; no hashes or signatures are included. Figure 4.4 illustrates this flow.




Figure 4.4: Data path from ingestion to curated output in the baseline pipeline.

4.4 Observed Gaps

The pipeline behaves correctly but lacks the mechanisms required to produce credible audit evidence. When compared against requirements R1–R5 from Chapter 3, several weaknesses become clear.

R1: Governance Accountability. Authorship and approval lineage are implicit. Commits identify contributors but are unsigned; Airflow and Terraform

run under shared service identities, leaving no trace between authorised decisions and executed actions. Accountability is assumed rather than enforced.

R2: Information Security Assurance. Infrastructure and data outputs are not verified for integrity. Containers and S3 objects are mutable; dependency scanning is absent. Runtime isolation exists, but cryptographic assurance does not. Evidence of reliability depends on logs that can be altered or lost.

R3: Operational Risk Control. Reproducibility relies on mutable Docker images and library versions. State files and metadata are not preserved, so rollback cannot be guaranteed. Identical code may yield divergent results under different environments, breaking referential consistency.

R4: Continuous Control Monitoring. The CI/CD workflow runs syntax checks but no policy validation. Separation of duties and multi-party approvals are not automated. Monitoring is reactive rather than preventative; control assurance cannot be measured continuously.

R5: End-to-End Evidence Chain. No unified record links input data, transformations, configurations, and outputs. Logs are scattered across components with no correlation or immutability. Lineage reconstruction requires manual collation, undermining the principles of referentiality and verifiability.

4.5 Summary

The baseline pipeline demonstrates a functioning but ungoverned ETL workflow. It moves data reliably from extraction to storage but depends on convention rather than control. Each tool—Airflow, PySpark, S3, Terraform, and GitHub Actions—fulfils its operational role but not its assurance role. None provide signed artifacts, immutable logs, or reproducible lineage. These gaps define the starting point for the DevSecOps architecture in Chapter 5, which integrates governance and verification to transform a working pipeline into an auditable one.

Chapter 5

Proposed Architecture

This chapter presents the proposed DevSecOps architecture for governing an ETL pipeline in a regulated environment. It builds directly on the baseline pipeline from Chapter 4, retaining the same essential data movement while adding the control, evidence, and verification mechanisms required to satisfy the governance objectives defined in Chapter 3.

The design should be read as a governance architecture rather than merely a runtime topology. Its purpose is not only to execute extraction, transformation, and loading tasks, but to ensure that each operational change, deployment, and dataset version can be explained after the fact using authoritative evidence. In this sense, the proposed contribution sits above the business logic of the pipeline and governs the conditions under which that logic is allowed to operate.

The architecture is described at escalating levels of detail. It begins with a lifecycle view of the governed system, then decomposes the major architectural subsystems, then narrows to the runtime data and metadata flows, and finally follows a representative run from change request to dataset version and audit reconstruction. Consistent with the methodology, the controls described here should be interpreted as representative minimum controls for the thesis prototype, not as a claim of production-complete assurance.

5.1 Lifecycle Overview

At the highest level, the proposed architecture links four concerns that are separate or weakly connected in the baseline: change intent, release control, runtime execution, and audit reconstruction. The central design claim is that a dataset should never appear as an isolated output. It should instead be the observable consequence of an approved change, an attested release, a specific infrastructure state, and a recorded execution context.

Figure 5.1 presents this lifecycle as a single governed process. The top lane captures change management, where a GitHub Issue establishes the initial statement of intent and a pull request subjects the change to automated and human

review. The second lane captures the execution plane, where an approved artifact is used to run the pipeline and produce a dataset version. The third lane introduces blocking controls, particularly artifact verification and runtime data quality gates. The fourth lane captures the evidence layer, where logs, run metadata, lineage metadata, and infrastructure references are retained so that an auditor can reconstruct what happened without manually correlating disconnected systems.




Figure 5.1: Unified lifecycle of the governed ETL architecture.

Change intent as the start of lineage. The evidence chain begins before the pipeline runs. A governed change must start with a documented issue or change request that records why the change is being made and what requirement or defect it addresses. This addresses Requirement R1 by preserving decision authority alongside code authorship.

Release control as the bridge between code and runtime. The architecture does not allow code to move directly from repository to execution environment by convention alone. It introduces CI/CD gates, policy evaluation, and artifact attestation so that runtime execution is tied to an approved release object rather than to an ambiguous repository state. This is the primary mechanism by which Requirements R2 and R4 are operationalised.

Runtime control as a governed execution path. The runtime path remains recognisably ETL: data are ingested, transformed, validated, and written to governed storage. What changes is that each step operates within a recorded control context. The run is identified, the deployed artifact is referenced, the infrastructure state is captured, and promotion to the final analytical layer is blocked unless quality conditions are met.

Audit reconstruction as an explicit design goal. The final concern is not throughput or convenience but reconstructability. Requirement R5 demands that

the system answer questions such as: Which change produced this dataset version? Which approved artifact ran? Which infrastructure state was active? Which control checks passed or failed? The architecture is therefore designed backwards from the audit query as much as it is designed forwards from the ETL process.

5.2 Structural Decomposition

Where the lifecycle view explains the system behaviourally, the high-level solution architecture explains it structurally. Figure 5.2 shows the major subsystems and boundaries that participate in the governed pipeline.




Figure 5.2: High-level solution architecture for the proposed pipeline.

External source boundary. The source system remains the SEC EDGAR API, preserving comparability with the baseline. This is important analytically: the proposed architecture is not an alternative business pipeline, but a governed form of the same core workload.

Delivery and control boundary. The first major subsystem is the delivery and control boundary. It combines the GitHub repository, pull request approval gate, GitHub Actions, Terraform, automated tests, security scanning, policy-as-code validation, and artifact attestation. This boundary is the control-plane entry point for all code and configuration change. In the implementation contract, it is also where the logical release manifest is assembled, linking the deployable Airflow and Glue assets to the CI run and attestation evidence that produced them.

AWS account boundary. Within the cloud target, the architecture is divided into three bands: *execution plane*, *data platform*, and *governance and evidence plane*. This separation matters because the thesis is evaluating more than runtime correctness. The execution plane runs the workflow, the data platform stores the outputs, and the governance plane preserves the context that makes those outputs auditable.

Execution plane. The execution plane consists of MWAA, AWS Glue Python jobs, and CloudWatch. MWAA acts as the orchestration surface for governed cloud execution, while Glue performs the operational work of ingestion, transformation, quality checking, and promotion. CloudWatch provides operational telemetry and supporting logs. During development, the same orchestration logic can be exercised locally in Docker-based Airflow, but the governed target architecture remains MWAA in the AWS environment.

Data platform. The data platform is intentionally layered. Raw extracted data are written to an S3 raw zone. Normalised and contract-coerced records are written to an S3 curated zone. Promoted analytical outputs are then written to the gold layer, represented conceptually as an S3 gold zone and technically as Amazon S3 Tables backed by Iceberg semantics. This progressive structure preserves the staging logic of ETL while distinguishing between transient intermediates and governed analytical state.

Governance and evidence plane. The governance and evidence plane contains Glue Crawler, Glue Data Catalog, OpenLineage, OpenMetadata, IAM Roles, and Secrets Manager. These components provide metadata context, service boundaries, secret handling, and queryable lineage. The design intent is that runtime metadata should not remain trapped inside operational logs. It should be surfaced into a metadata plane that can join lineage events, catalog state, change references, and artifact provenance.

Downstream consumption. The architecture ends with two downstream actors: the analytics consumer and the audit or lineage query user. This distinction is deliberate. The analytics consumer needs governed data products; the auditor

needs reconstructed evidence. A mature architecture must satisfy both, but they query different parts of the system for different reasons.

5.3 Governed Process Maps

The high-level structure becomes clearer when decomposed into the concrete processes that operate across it. The process maps in this section show how the same architecture behaves under different concerns: change governance, continuous delivery, infrastructure change, runtime execution, and audit reconstruction.

5.3.1 Change Management Workflow

Figure 5.3 shows how the architecture treats change management as the first control surface rather than as an informal project-management layer.




Figure 5.3: Governed change management workflow.

The process begins with a change request and a GitHub Issue containing design notes and acceptance criteria. Development then proceeds on a branch that references the ticket identifier, ensuring that the repository history carries forward the originating governance reference. A pull request triggers automated checks covering linting, tests, policy validation, security review, and ticket enforcement. If those checks fail, the change returns for remediation; if they pass, the change still requires explicit human approval from the Code Owner and Data Owner before merge.

This workflow operationalises Requirement R1 by binding intent, authorship, and approval into a single continuous chain. It also contributes to Requirement R4, because policy and validation are not optional review activities but blocking controls in the merge path.

5.3.2 CI/CD Release Workflow

Figure 5.4 narrows to the release path itself.




Figure 5.4: CI/CD workflow for packaging, attestation, and deployment.

Once a change is merged, GitHub Actions performs build and lint checks, unit and integration testing, security scanning, and policy-as-code validation. Only then is the deployment artifact built, attested, and released for deployment. This sequencing is critical. In the baseline, deployment is driven primarily by repository events; in the proposed architecture, deployment is driven by the successful creation of an approved release object. The conceptual trace becomes: `commit -> CI run -> approved artifact -> attestation -> release -> deployment`.

This is the principal mechanism for satisfying Requirement R2. Integrity is not assumed because a file exists in the repository; it is asserted only after the artifact has passed the required gates and has provenance attached to it.

5.3.3 Infrastructure Change Workflow

Figure 5.5 shows that infrastructure change is governed using the same logic as application change.




Figure 5.5: Infrastructure change workflow with Terraform state capture.

Terraform changes begin with an issue, pass through pull-request review, and are subjected to validation and policy checks before merge. After `terraform apply`, the applied infrastructure state version is recorded as evidence. This gives the architecture a durable reference to the platform configuration that existed at the time a dataset was produced. That reference is essential for Requirement R3 because reproducibility depends not only on code version, but also on the state of the environment in which the code ran.

5.3.4 Runtime ETL Execution Workflow

Figure 5.6 shows the governed runtime path.




Figure 5.6: Governed runtime execution workflow.

The flow begins with a pipeline trigger and Airflow DAG creation, followed by verification that the runtime is using an approved artifact. Glue jobs then read the source dataset, perform schema validation, run data quality and data contract checks, apply the transformation, and produce curated output. Only after the post-transform checks pass is the dataset promoted to the gold layer. Execution logs, run evidence, and lineage evidence are recorded as separate but linked outputs.

The important difference from the baseline is that quality and promotion are not descriptive afterthoughts. They are blocking decisions in the runtime path. This converts data quality from an operational best practice into a control instrument that directly governs whether an output is permitted to become a governed dataset version.

5.3.5 Lineage and Audit Reconstruction Workflow

Figure 5.7 reverses the direction of reasoning and begins with the audit question instead of the pipeline trigger.




Figure 5.7: Audit reconstruction workflow from dataset to change intent.

An audit request begins with a dataset under investigation. From that gold dataset, the system performs a lineage lookup to identify the pipeline run ID, retrieves runtime metadata as authoritative evidence, consults execution logs as supporting evidence, and identifies the artifact version, Git commit, infrastructure state version, and originating change request. The intended result is a coherent reconstruction of **data + code + infrastructure + intent**.

This workflow directly expresses Requirement R5. A system that can only produce logs has not solved the governance problem; it has only moved the manual work to a later stage. The proposed design aims to make reconstruction a first-class system capability.

5.4 Data and Metadata Flows

The process maps describe behaviour over time. The data-flow diagrams explain how runtime assets, datasets, and evidence objects move between components. Three views are important: the overall system data flow, the governed runtime flow, and the metadata and evidence flow.

5.4.1 System Data Flow

Figure 5.8 presents the end-to-end system data flow across the repository, delivery controls, AWS runtime, metadata services, and downstream consumers.




Figure 5.8: System-level data flow across delivery, runtime, and metadata layers.

The governing repository and CI/CD systems produce an approved workflow artifact. That artifact, together with Terraform-provisioned infrastructure, establishes the approved runtime context for MWAA and Glue. Runtime data then move from SEC EDGAR into the raw zone, from raw to curated, and from curated to the gold layer. The gold layer is surfaced as governed tables through S3 Tables/Iceberg, with Glue Crawler and Glue Data Catalog providing additional metadata context.

This system view shows why the proposed architecture is more than an orches-

tration pattern. The delivery path, runtime path, and metadata path are linked so that the data product and the control product remain synchronised.

5.4.2 Runtime Data Flow

Figure 5.9 narrows from the whole system to a single governed execution context.




Figure 5.9: Runtime flow showing the approved execution context and control gates.

This view emphasises that runtime execution is dependent on prior control outcomes. External triggering alone is insufficient. The run must also inherit an

approved workflow artifact, its attestation, the relevant Terraform state version, and the evidence that required policy and security gates have already passed. Once in runtime, the flow proceeds through MWAA, Glue ingestion, schema validation, the raw zone, Glue transformation, the curated zone, data quality checks, data contract validation, the promotion gate, and finally the gold dataset and table abstraction.

The diagram therefore combines two forms of control. The first is pre-execution control, inherited from CI/CD. The second is in-execution control, applied to data and transformation results before promotion. Together they bind Requirements R2, R3, and R4 to the same execution path.

5.4.3 Metadata and Evidence Flow

Figure 5.10 describes the movement of evidence and references rather than business data.




Figure 5.10: Metadata and evidence flow supporting audit reconstruction.

The metadata flow begins in change management with the issue and pull request, continues through GitHub Actions, OPA policy evaluation, attestation, approved artifact version, Terraform, and infrastructure state version, then enters the runtime as MWWA and Glue run metadata. At runtime, dataset versions, lineage events, run records, and catalog metadata are pushed into OpenMeta-

data, which acts as the query hub for reconstructed evidence chains. The intended audit path is explicit in the diagram: `dataset -> run -> artifact -> infrastructure -> change`.

This design choice is central to the thesis. Logs alone are retained only as supporting evidence. The authoritative audit chain is constructed from linked metadata references that can be queried consistently and validated independently.

5.5 Worked Example: One Governed EDGAR Run

To make the architecture concrete, this section follows the representative pipeline slice implemented in the prototype repository. The current codebase models a single governed EDGAR flow centred on the `edgar_governed_pipeline`, with separate ingestion, transformation, quality, and promotion steps plus a versioned run manifest.

1. Change and release preparation. A change begins as a GitHub Issue and is implemented on a branch that encodes the governing issue identifier. The pull request body carries phase and requirement references, which are parsed into machine-readable governance inputs for CI policy checks. Automated validation then evaluates branch naming, issue linkage, template completeness, tests, and policy conditions before the change is allowed to merge. Once approved, the release process produces a logical release object that binds together the workflow artifact, the CI run, the attestation record, and the deployment target.

2. Run initialisation. When the DAG is triggered, it first creates a run context and an initial manifest. In the current prototype, that manifest records the `run_id`, `pipeline_id`, `dataset_date`, `source_ref`, `release_manifest_ref`, `terraform_state_ref`, `change_ref`, the expected job references, the gold table coordinates, and the initial status. This occurs before any data movement, which means the execution is evidence-bearing from its first step rather than only at completion.

3. Ingestion. The ingestion job reads the EDGAR source payload, expands the `filings.recent` structure into row-wise records, normalises the key fields, and writes two raw outputs: the original payload and the normalised record set. In the current implementation these outputs are partitioned by `dataset_date` and `run_id`. The manifest is then updated with the raw output locations and the raw row count. At this point, the run can already prove which source was consulted and where the first persisted outputs were written.

4. Transformation and contract coercion. The transformation job reads the normalised raw records, enriches them with the run identifier and dataset date, coerces them to the curated data contract, and writes the curated dataset to the curated zone. The manifest is updated again with the curated output location, the contract reference, and the curated row count. This step is significant because it converts the raw landing zone into a schema-governed intermediate suitable for policy-based promotion.

5. Runtime quality gate. The data quality step evaluates the curated output against a small but explicit control set: the curated object must exist, it must match the declared contract, the row count must be greater than zero, and required identifiers must be populated. The result is recorded in the manifest as a structured `quality_result`. If any check fails, the run is finalised with a `failed_quality` status and no dataset version is emitted. This is the runtime expression of the blocking controls shown in the process maps.

6. Promotion and dataset versioning. If the quality gate passes, the promotion step writes the approved records to the gold table target. The current prototype computes a deterministic dataset version from the gold table identity and a digest of the promoted content, then stores both the dataset version and the write result in the manifest. This matters for Requirement R3: replayability and reconstruction are only credible if the final version identifier is stable for identical governed inputs and execution conditions.

7. Final manifest and audit lookup. At the end of the run, the manifest contains the raw outputs, curated outputs, quality result, row counts, dataset version, gold write result, timestamps, and final status. In the integration flow used by the prototype, a representative run is linked to references such as `release-manifests/dev-edgar-v1.js`, `terraform-state/dev/serial-0001`, and `issue/EDGAR-1`. An auditor starting from the resulting dataset version can therefore traverse back to the run manifest, then to the release manifest, infrastructure state reference, and governing change request. The evidence chain is not inferred; it is recorded explicitly.

Worked-example significance. This example illustrates the architectural argument of the thesis. The novelty is not that EDGAR data can be ingested or transformed; the baseline already demonstrates that. The novelty is that the governed run emits enough linked context to explain why the dataset exists, under what authority it was produced, which control conditions were satisfied, and how the same state could be replayed or challenged later.

5.6 Summary

The proposed architecture transforms the baseline ETL pipeline into a governed data-delivery system. It introduces a control plane that begins with change intent, continues through automated and human release gates, governs runtime promotion decisions, and preserves a queryable evidence chain spanning code, configuration, data, and infrastructure.

At the conceptual level, the architecture integrates change management, continuous delivery, runtime orchestration, metadata governance, and audit reconstruction into one lifecycle. At the structural level, it separates the execution plane, data platform, and governance plane while linking them through attested artifacts, infrastructure references, lineage events, and run manifests. At the operational level, it culminates in a governed pipeline run whose dataset version can be reconstructed back to its originating change request. This architecture provides the design foundation for the prototype and the evaluation that follow.

Chapter 6

Prototype Implementation

To be completed in Thesis B

Chapter 7

Evaluation

To be completed in Thesis B

Chapter 8

Discussion

To be completed in Thesis B

Chapter 9

Conclusion

To be completed in Thesis B

Appendix A

Abbreviations

APRA	Australian Prudential Regulation Authority
AWS	Amazon Web Services
CI/CD	Continuous Integration / Continuous Delivery
CPG	Prudential Practice Guide (APRA)
CPS	Prudential Standard (APRA)
DAST	Dynamic Application Security Testing
DAG	Directed Acyclic Graph
DLC	Data Lifecycle
DSR	Design Science Research
ELT	Extract–Load–Transform
EMR	Elastic MapReduce
EC2	Elastic Compute Cloud
ETL	Extract–Transform–Load
GDPR	General Data Protection Regulation
GHCR	GitHub Container Registry
GitOps	Git-based Operations
HIPAA	Health Insurance Portability and Accountability Act
IaC	Infrastructure as Code
IAM	Identity and Access Management
IEEE	Institute of Electrical and Electronics Engineers
ISO	International Organization for Standardization
ML	Machine Learning
MLOps	Machine Learning Operations
NIST	National Institute of Standards and Technology
PaC	Policy as Code
RBAC	Role-Based Access Control
S3	Simple Storage Service
SAST	Static Application Security Testing
SCA	Software Composition Analysis
SBOM	Software Bill of Materials
SDLC	Software Development Lifecycle

SSDF Secure Software Development Framework

Bibliography

- [1] “ISO/IEC/IEEE International Standard – Systems and software engineering - Content of life-cycle information items (documentation),” Jul. 2019. [Online]. Available: <https://ieeexplore.ieee.org/document/8767110>
- [2] B. Achuthan and M. A. Alimohideen, “Shifting Gears: Integrating Security Audits into Automotive DevSecOps,” in *2024 International Conference on Vehicular Technology and Transportation Systems (ICVTTS)*, vol. 1, Sep. 2024, pp. 1–6. [Online]. Available: <https://ieeexplore.ieee.org/document/10763937>
- [3] M. Anisetti, C. A. Ardagna, E. Damiani, and F. Gaudenzi, “A Semi-Automatic and Trustworthy Scheme for Continuous Cloud Service Certification,” *IEEE Transactions on Services Computing*, vol. 13, no. 1, pp. 30–43, Jan. 2020. [Online]. Available: <https://ieeexplore.ieee.org/abstract/document/7831357/authors>
- [4] M. Anisetti, N. Bena, F. Berto, and G. Jeon, “A DevSecOps-based Assurance Process for Big Data Analytics,” in *2022 IEEE International Conference on Web Services (ICWS)*, Jul. 2022, pp. 1–10. [Online]. Available: <https://ieeexplore.ieee.org/abstract/document/9885738>
- [5] Australian Prudential Regulation Authority, “Prudential Practice Guide CPG 235: Managing Data Risk,” Australian Prudential Regulation Authority (APRA), Canberra, ACT, Prudential Practice Guide CPG 235, Sep. 2013. [Online]. Available: https://www.apra.gov.au/sites/default/files/Prudential-Practice-Guide-CPG-235-Managing-Data-Risk_1.pdf
- [6] —, “Prudential Standard CPS 234: Information Security,” Canberra, ACT, Jul. 2019. [Online]. Available: https://www.apra.gov.au/sites/default/files/cps_234_july_2019_for_public_release.pdf
- [7] —, “Prudential Standard CPS 510: Governance,” Canberra, ACT, Jul. 2023. [Online]. Available: https://www.apra.gov.au/sites/default/files/prudential_standard_cps.510_governance.pdf
- [8] —, “Prudential Standard CPS 230: Operational Risk Management,” Canberra, ACT, Jul. 2025. [Online]. Available: https://www.apra.gov.au/sites/default/files/cps_230_july_2025_for_public_release.pdf

- [//www.apra.gov.au/sites/default/files/2025-07/Prudential%20Standard%20CPS%20230%20Operational%20Risk%20Management%20-%20clean.pdf](https://www.apra.gov.au/sites/default/files/2025-07/Prudential%20Standard%20CPS%20230%20Operational%20Risk%20Management%20-%20clean.pdf)
- [9] A. Ball, *Review of Data Management Lifecycle Models*. Bath, UK: University of Bath, Feb. 2012.
- [10] C. Castellanos, C. A. Varela, and D. Correal, “ACCORDANT: A domain specific-model and DevOps approach for big data analytics architectures,” *Journal of Systems and Software*, vol. 172, p. 110869, Feb. 2021. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0164121220302594>
- [11] T. Chen and H. Suo, “Design and Practice of Security Architecture via DevSecOps Technology,” in *2022 IEEE 13th International Conference on Software Engineering and Service Science (ICSESS)*, Oct. 2022, pp. 310–313, ISSN: 2327-0594. [Online]. Available: <https://ieeexplore.ieee.org/document/9930212>
- [12] D. B. Cruz, J. R. Almeida, and J. L. Oliveira, “Open Source Solutions for Vulnerability Assessment: A Comparative Analysis,” *IEEE Access*, vol. 11, pp. 100 234–100 255, 2023. [Online]. Available: <https://ieeexplore.ieee.org/document/10251527>
- [13] P. Cupoli, S. Earley, and D. Henderson, “DAMA-DMBOK2 Framework.”
- [14] S. Gregor and A. R. Hevner, “Positioning and Presenting Design Science Research for Maximum Impact,” *MIS Quarterly*, vol. 37, no. 2, pp. 337–355, 2013, publisher: Management Information Systems Research Center, University of Minnesota. [Online]. Available: <https://www.jstor.org/stable/43825912>
- [15] S. Gupta, M. Bhatia, M. Memoria, and P. Manani, “Prevalence of GitOps, DevOps in Fast CI/CD Cycles,” in *2022 International Conference on Machine Learning, Big Data, Cloud and Parallel Computing (COM-IT-CON)*, vol. 1, May 2022, pp. 589–596. [Online]. Available: <https://ieeexplore.ieee.org/document/9850786>
- [16] R. Hernández, B. Moros, and J. Nicolás, “Requirements management in DevOps environments: a multivocal mapping study,” *Requirements Engineering*, vol. 28, no. 3, pp. 317–346, Sep. 2023. [Online]. Available: <https://doi.org/10.1007/s00766-023-00396-w>
- [17] A. Hevner, A. R. S. March, S. T. Park, J. Park, Ram, and Sudha, “Design Science in Information Systems Research,” *Management Information Systems Quarterly*, vol. 28, pp. 75–, Mar. 2004.

- [18] R. Kimball and J. Caserta, *The Data Warehouse ETL Toolkit: Practical Techniques for Extracting, Cleaning, Conforming, and Delivering Data*. Wiley, 2004, iSBN: 9780764567575. [Online]. Available: <https://www.oreilly.com/library/view/the-data-warehouse/9780764567575/>
- [19] B. O. Kose, “Agility Meets Compliance: The Convergence of Data Governance and DevSecOps,” in *Data Governance, DevSecOps, and Advancements in Modern Software*, B. E. Elbaghazaoui, M. Amnai, and N. Gherabi, Eds. IGI Global, Apr. 2025, pp. 131–154. [Online]. Available: <https://services.igi-global.com/resolvedoi/resolve.aspx?doi=10.4018/979-8-3373-0365-9.ch007>
- [20] W. C. Lenhardt, S. Ahalt, B. Blanton, L. Christopherson, and R. Idaszak, “Data Management Lifecycle and Software Lifecycle Management in the Context of Conducting Science | Journal of Open Research Software,” Feb. 2014. [Online]. Available: <https://openresearchsoftware.metajnl.com/articles/10.5334/jors.ax>
- [21] J. Li and H. Li, “Evolution of Application Security based on OWASP Top 10 and CWE/SANS Top 25 with Predictions for the 2025 OWASP Top 10,” in *2025 International Conference on Inventive Computation Technologies (ICICT)*, Apr. 2025, pp. 1178–1183, iSSN: 2767-7788. [Online]. Available: <https://ieeexplore.ieee.org/document/11004742>
- [22] D. Mahmud and M. Z. Ikbali, “THE ROLE OF ETL (EXTRACT-TRANSFORM-LOAD) PIPELINES IN SCALABLE BUSINESS INTELLIGENCE: A COMPARATIVE STUDY OF DATA INTEGRATION TOOLS,” *ASRC Procedia: Global Perspectives in Science and Scholarship*, vol. 2, no. 1, pp. 89–121, Apr. 2022. [Online]. Available: <https://global.asrcconference.com/index.php/asrc/article/view/29>
- [23] A. Mittal and V. Venkatesan, “Practical Integration of Large Language Models into Enterprise CI/CD Pipelines for Security Policy Validation: An Industry-Focused Evaluation,” in *2025 IEEE International Conference on Service-Oriented System Engineering (SOSE)*, Jul. 2025, pp. 197–203, iSSN: 2642-6587. [Online]. Available: <https://ieeexplore.ieee.org/document/11126149>
- [24] H. A. Mohna, T. Barua, M. Mohiuddin, and M. M. Rahman, “AI-READY DATA ENGINEERING PIPELINES: A REVIEW OF MEDALLION ARCHITECTURE AND CLOUD-BASED INTEGRATION MODELS,” *American Journal of Scholarly Research and Innovation*, vol. 1, no. 01, pp. 319–350, Apr. 2022. [Online]. Available: <https://researchinnovationjournal.com/index.php/AJSRI/article/view/51>

- [25] A. R. Munappy, D. I. Mattos, J. Bosch, H. H. Olsson, and A. Dakkak, "From Ad-Hoc Data Analytics to DataOps," in *Proceedings of the International Conference on Software and System Processes*, ser. ICSSP '20. New York, NY, USA: Association for Computing Machinery, 2020, pp. 165–174, event-place: Seoul, Republic of Korea. [Online]. Available: <https://doi.org/10.1145/3379177.3388909>
- [26] H. Myrbakken and R. Colomo-Palacios, "DevSecOps: A Multivocal Literature Review," in *Software Process Improvement and Capability Determination*, A. Mas, A. Mesquida, R. V. O'Connor, T. Rout, and A. Dorling, Eds. Cham: Springer International Publishing, 2017, pp. 17–29.
- [27] S. Nai, A. Rifai, and A. Sadiq, "Data Governance, Key Insights, Strategic Challenges, and Future Imperatives," in *Data Governance, DevSecOps, and Advancements in Modern Software*, B. E. Elbaghazaoui, M. Amnai, and N. Gherabi, Eds. IGI Global, Apr. 2025, pp. 1–16. [Online]. Available: <https://services.igi-global.com/resolvedoi/resolve.aspx?doi=10.4018/979-8-3373-0365-9.ch001>
- [28] P. Pandey and A. Patel, "Integrating Security in Cloud-Native Development: A DevSecOps Approach to Resilient Software Systems," in *Data Governance, DevSecOps, and Advancements in Modern Software*, B. E. Elbaghazaoui, M. Amnai, and N. Gherabi, Eds. IGI Global, Apr. 2025, pp. 169–196. [Online]. Available: <https://services.igi-global.com/resolvedoi/resolve.aspx?doi=10.4018/979-8-3373-0365-9.ch009>
- [29] A. Patel and P. Pandey, "Safeguarding Sensitive Data in the DevSecOps Pipelines: Strategies, Tools, and Best Practices for Modern Software Development," in *Data Governance, DevSecOps, and Advancements in Modern Software*, B. E. Elbaghazaoui, M. Amnai, and N. Gherabi, Eds. IGI Global, Apr. 2025, pp. 215–240. [Online]. Available: <https://services.igi-global.com/resolvedoi/resolve.aspx?doi=10.4018/979-8-3373-0365-9.ch011>
- [30] K. Peffers, T. Tuunanen, M. A. Rothenberger, and S. Chatterjee, "A Design Science Research Methodology for Information Systems Research," *Journal of Management Information Systems*, vol. 24, no. 3, pp. 45–77, 2007, publisher: Taylor & Francis, Ltd. [Online]. Available: <https://www.jstor.org/stable/40398896>
- [31] A. Pogiatis and G. Samakovitis, "An Event-Driven Serverless ETL Pipeline on AWS," *Applied Sciences*, vol. 11, no. 1, p. 191, Jan. 2021, publisher: Multidisciplinary Digital Publishing Institute. [Online]. Available: <https://www.mdpi.com/2076-3417/11/1/191>
- [32] T. Rangnau, R. v. Buijtenen, F. Fransen, and F. Turkmen, "Continuous Security Testing: A Case Study on Integrating Dynamic Security Testing

- Tools in CI/CD Pipelines,” in *2020 IEEE 24th International Enterprise Distributed Object Computing Conference (EDOC)*, Oct. 2020, pp. 145–154, iSSN: 2325-6362. [Online]. Available: <https://ieeexplore.ieee.org/abstract/document/9233212>
- [33] M. Souppaya, K. Scarfone, and D. Dodson, “Secure Software Development Framework (SSDF) Version 1.1: Recommendations for Mitigating the Risk of Software Vulnerabilities,” National Institute of Standards and Technology, Tech. Rep. NIST Special Publication (SP) 800-218, Feb. 2022. [Online]. Available: <https://csrc.nist.gov/pubs/sp/800/218/final>
- [34] K. Soussane, B. E. Elbaghazaoui, and M. Amnai, “Integrating Security in DataOps: A Framework for Securing Data Pipelines in Real-time Analytics,” in *Data Governance, DevSecOps, and Advancements in Modern Software*, B. E. Elbaghazaoui, M. Amnai, and N. Gherabi, Eds. IGI Global, Apr. 2025, pp. 197–214. [Online]. Available: <https://services.igi-global.com/resolvedoi/resolve.aspx?doi=10.4018/979-8-3373-0365-9.ch010>
- [35] H. V. Sreepathy, B. Dinesh Rao, M. Kumar Jaysubramanian, and B. Deepak Rao, “Data Ingestions as a Service (DIaaS): A Unified Interface for Heterogeneous Data Ingestion, Transformation, and Metadata Management for Data Lake,” *IEEE Access*, vol. 12, pp. 156 131–156 145, 2024. [Online]. Available: <https://ieeexplore.ieee.org/document/10716380/>
- [36] R. Vayyala, “Data Lineage: Tracing Data Across Systems,” in *Data Governance, DevSecOps, and Advancements in Modern Software*, B. E. Elbaghazaoui, M. Amnai, and N. Gherabi, Eds. IGI Global, Apr. 2025, pp. 73–98. [Online]. Available: <https://services.igi-global.com/resolvedoi/resolve.aspx?doi=10.4018/979-8-3373-0365-9.ch004>
- [37] —, “Ensuring Data Quality and Integrity in Modern Enterprises:,” in *Data Governance, DevSecOps, and Advancements in Modern Software*, B. E. Elbaghazaoui, M. Amnai, and N. Gherabi, Eds. IGI Global, Apr. 2025, pp. 17–46. [Online]. Available: <https://services.igi-global.com/resolvedoi/resolve.aspx?doi=10.4018/979-8-3373-0365-9.ch002>
- [38] —, “Metadata Management and Its Role in Data Governance:,” in *Data Governance, DevSecOps, and Advancements in Modern Software*, B. E. Elbaghazaoui, M. Amnai, and N. Gherabi, Eds. IGI Global, Apr. 2025, pp. 47–72. [Online]. Available: <https://services.igi-global.com/resolvedoi/resolve.aspx?doi=10.4018/979-8-3373-0365-9.ch003>
- [39] Q. Wang, C. Wang, K. Ren, W. Lou, and J. Li, “Enabling Public Auditability and Data Dynamics for Storage Security in Cloud Computing,” *IEEE Transactions on Parallel and Distributed Systems*, vol. 22, no. 5, pp. 847–859, May 2011. [Online]. Available: <https://ieeexplore.ieee.org/document/5611497>

- [40] H. Weigand and P. Elsas, “Auditability as a Design Problem,” in *2019 IEEE 21st Conference on Business Informatics (CBI)*, vol. 01, Jul. 2019, pp. 276–285, iSSN: 2378-1971. [Online]. Available: <https://ieeexplore.ieee.org/abstract/document/8808092>
- [41] J. M. Wing, “The Data Life Cycle,” *Harvard Data Science Review*, vol. 1, no. 1, Jul. 2019, publisher: The MIT Press. [Online]. Available: <https://hdsr.mitpress.mit.edu/pub/577rq08d/release/4>
- [42] W. Yang, R. Fu, M. B. Amin, and B. Kang, “The Impact of Modern AI in Metadata Management,” *Human-Centric Intelligent Systems*, vol. 5, no. 3, pp. 323–350, Sep. 2025. [Online]. Available: <https://doi.org/10.1007/s44230-025-00106-5>
- [43] J. Zhang, J. Symons, P. Agapow, J. T. Teo, C. A. Paxton, J. Abdi, H. Mattie, C. Davie, A. Z. Torres, A. Folarin, H. Sood, L. A. Celi, J. Halamka, S. Eapen, and S. Budhdeo, “Best practices in the real-world data life cycle,” *PLOS Digital Health*, vol. 1, no. 1, p. e0000003, Jan. 2022, publisher: Public Library of Science. [Online]. Available: <https://journals.plos.org/digitalhealth/article?id=10.1371/journal.pdig.0000003>
- [44] Y. Zhou, Y. Yu, and B. Ding, “Towards MLOps: A Case Study of ML Pipeline Platform,” in *2020 International Conference on Artificial Intelligence and Computer Engineering (ICAICE)*, Oct. 2020, pp. 494–500. [Online]. Available: <https://ieeexplore.ieee.org/abstract/document/9361315>